

First Best Implementation with Costly Information Acquisition^{*}

Daniil Larionov[†] Hien Pham[‡] Takuro Yamashita[§] Shuguang Zhu[¶]

First version: October 14, 2021

This version: November 15, 2022

Click [here](#) for the latest version.

Abstract

We study mechanism design with flexible but costly information acquisition. There is a principal and four or more agents, sharing a common prior over the set of payoff-relevant states. The principal proposes a mechanism to the agents who can then acquire information about the state of the world by privately designing a signal device. As long as it is costless for each agent to acquire a signal that is independent from the state, we show that there exists a mechanism which allows the principal to implement any social choice rule at zero information acquisition cost to the agents.

^{*}Daniil Larionov acknowledges funding by the German Research Foundation (DFG) through CRC TR 224 (Project B01). Takuro Yamashita acknowledges funding from the European Research Council (ERC) under the European Union’s Horizon 2020 research and innovation program (grant agreement No 714693), and ANR under grant ANR-17-EURE-0010 (Investissements d’Avenir program). Shuguang Zhu is supported by “the Fundamental Research Funds for the Central Universities”.

[†]ZEW Mannheim. daniil.larionov@zew.de

[‡]Toulouse School of Economics, University of Toulouse Capitole. hien.pham@tse-fr.eu

[§]Osaka School of International Public Policy, Osaka University. yamashita.takuro.osipp@osaka-u.ac.jp

[¶]School of Economics, Shanghai University of Finance and Economics. zhu.shuguang@mail.shufe.edu.cn

1 Introduction

In most mechanism design problems, there is a collection of agents who have exogenously given private information, and there is a principal who desires to implement a social choice rule by designing a mechanism which incentivizes the agents to reveal their information.

In many practical problems, however, the agents' private information is often a consequence of their own (possibly costly) *information acquisition*. For example, bidders in an oil-tract auction (Wilson, 1969) may conduct test drills; bidders in a spectrum auction may conduct market research; voters in a presidential election may investigate the candidates' past political activities; members of a hiring committee may study the job applicant's background in order to see whether he is fit for the job.

Importantly, in such situations, a mechanism in place does not only affect each agent's incentive of reporting the acquired information truthfully, but also affects his choice of *what kind of information* to acquire. In this sense, the properties of desirable mechanisms could potentially be very different from those which only guarantee truth-telling incentives for a *given* information structure.

Although this issue is already relevant in single-agent environments,¹ the degree of complexity is even higher in multi-agent environments: in principle, flexibility of each agent's information acquisition does not only mean flexibility in terms of his signal's informativeness about the payoff-relevant state, but also means flexibility in terms of his signal's informativeness about his opponents' signals. This issue of higher-order information and beliefs distinguishes multi-agent from single-agent environments. Modeling the dependence of the cost of information acquisition on higher-order information is a challenging task. In this paper, we assume that an agent's information acquisition cost only depends on his signal's informativeness about the payoff-relevant state, but not about the other agents' signals. In particular, it is costless to acquire a signal that is independent from the payoff-relevant state.² For example, imagine the situation where agents (e.g., telecom companies who buy spectrum) have to acquire information

¹Mensch (2020) studies a mechanism design problem with a single agent. See also Section 1.1.

²The literature on cost of information proposes and discusses a variety of possible cost functions (see Section 1.1), but it seems to be universally accepted that uninformative signals about the state of the world are costless.

from data providers (e.g., market research firms) operating on a competitive market for data. Each data provider generates signals about the payoff-relevant state (e.g., demand conditions in the mobile services market). Competition among the data providers forces them to price their data at the cost of production, which in turn depends on the informativeness of their data. The agents can then decide to make their signals perfectly correlated by strategically choosing the same data provider, or less than perfectly correlated by choosing different data providers. In both cases, the agents will pay the same price for the same informativeness, and hence the cost of information acquisition will be independent of the correlation structure among signals.

Of course, one can come up with cases where other cost specifications seem more reasonable (e.g., more positive correlation costlier, or less costlier: We discuss a range of possible alternative assumptions in Section 5.1). However, we think that this is a reasonable assumption as the first step toward the analysis of mutual information acquisition environments.

We consider a model with four or more agents. The principal and the agents share a common prior about the payoff-relevant state, and none of them has any private information at the beginning. We show that there exists a mechanism which allows the principal to implement any social choice rule at zero information acquisition cost to the agents. The key idea is that the mechanism recommends each agent to choose a special information acquisition action, which satisfies the *individually-uninformative-but-aggregately-revealing* property of [Zhu \(2021\)](#) (and each agent finds it optimal to obey this recommended action). The *individually-uninformative* part means that each agent's signal on its own is independent from the payoff-relevant state, which guarantees that his information cost is zero. The *aggregately-revealing* part means that the principal, by observing all the agents' reports — in fact, any two of them — can correctly identify the true payoff-relevant state. The fact that only two are enough, together with the fact that there are four or more agents, enables the principal to detect any unilateral deviation. It thus establishes the incentive compatibility of the mechanism.

1.1 Related Literature

In the literature on information acquisition in mechanism design, we usually consider restricted and/or less flexible spaces of information (see, for example, [Bergemann and Välimäki \(2002\)](#) for efficient mechanism design, [Shi \(2012\)](#) and [Bikhchandani and Obara \(2017\)](#) for optimal auction design, and [Persico \(2004\)](#), [Gerardi and Yariv \(2008\)](#), [Gershkov and Szentes \(2009\)](#), and [Zhao \(2016\)](#) for committee design with information acquisition.³)

[Mensch \(2020\)](#) studies mechanism design with a single agent’s flexible and costly information acquisition, building on the rational inattention framework ([Sims \(2003\)](#))⁴. Flexible and costly information acquisition is also considered by [Gleyze and Pernoud \(2020\)](#) who study a mechanism design problem with transferable utility and private values, in which agents acquire costly information on their own preferences and the preferences of other agents, and by [Ravid et al. \(2020\)](#) who study a bilateral trade model with costly information acquisition by the buyer. Flexible but not costly information acquisition is considered by [Roesler and Szentes \(2017\)](#) in the context of buyer-optimal information in monopoly pricing,⁵ by [Bergemann et al. \(2017\)](#) and [Brooks and Du \(2021\)](#) in the context of seller-pessimal information in common-value auctions, and by [Yamashita \(2018\)](#) in private-value auctions. All these papers feature a single entity, “nature”, who chooses the information structure (of one or multiple agents). In contrast to that, in our model each agent acquires information in a decentralized manner, which leads to a very different conclusion.

The information structure we employ was proposed in the context of mechanism design by [Zhu \(2021\)](#), who studies information disclosure by a mechanism designer. It builds on the idea of the one-time pad, an unbreakable encryption method ([Shannon, 1949](#)).⁶

³Restricting to the class of conservative rules, [Li \(2001\)](#) solves for the optimal degree of conservatism in committee design. The optimally chosen conservative rule outperforms the *ex post* optimal rule.

⁴[Mensch \(2020\)](#) also considers a multiple-agent extension of his model, but restricts attention to symmetric mechanisms in an independent private values setting, in which agents can acquire information about their own values, but cannot acquire any information about others’ values.

⁵See also [Condorelli and Szentes \(2020\)](#), though they also consider non-information changes of the agent’s private information distribution.

⁶See also [Krähmer \(2020\)](#) and [Krähmer \(2021\)](#) in the context of information disclosure in mechanism design

This key information structure makes the agents’ acquired information statistically dependent. In quasi-linear environments, [Cr  mer and McLean \(1988\)](#) show that the principal can extract full surplus from the agents who share a correlated prior. Although the extreme positivity of the results is a common feature of our paper and theirs, the two problems are quite different. First, our paper does not assume quasi-linearity. Second, their side-bet mechanism exploits an exogenously given correlated signal structure, and it is not clear if such a signal structure can be induced in equilibrium given some reasonable space of information acquisition actions.⁷ In our case, the resulting information structure is an equilibrium outcome, even though each agent can potentially acquire information independently from the others’ signals.

In non-quasi-linear environments, such as collective decision-making in committees, the first best outcome is generally not implementable under the commonly imposed restrictions on information acquisition technologies. For example, [Li \(2001\)](#) and [Persico \(2004\)](#), assuming that the agents have access to conditionally independent signals, show that the first best outcome is not attainable. In contrast to the previous results, we show that correlated information acquisition helps to implement the first best outcome.

There is a growing literature on the cost of flexible information in decision environments (see for example [Sims \(2003\)](#), [Matejka and McKay \(2015\)](#), [Caplin and Dean \(2015\)](#), and [Pomatto et al. \(2020\)](#)). Usually the main focus is on the cost of acquiring more or less precise information about a payoff-relevant state, and its relationship with a single decision-maker’s optimal choice. The framework, however, has been applied in multi-player problems, e.g. in coordination games ([Yang \(2015\)](#); [Morris and Yang \(2021\)](#); [Denti \(2020\)](#)). In particular, [Denti \(2020\)](#) proposes a model of unrestricted information acquisition in games, in which, as in our paper, the players can endogenously learn about a payoff-relevant state and actions of other players.

and strategic communication respectively where the randomization of information structures is allowed to keep the single agent (sender) uninformative; [Kalai et al. \(2010\)](#), [Renou and Tomala \(2012\)](#), [Renault et al. \(2014\)](#) in the context of games of communication network. [Peters and Troncoso-Valverde \(2013\)](#) apply this idea in mechanism-design games with multiple principals, and [Liu \(2015\)](#) applies it in his concept of individually uninformative correlating device. Our construction is most directly related to [Zhu \(2021\)](#).

⁷[Bikhchandani \(2010\)](#) shows that, indeed, an agent in the Cr  mer-McLean mechanism may have a strong incentive of acquiring information about others.

2 Model

2.1 Setup

There is a principal and $I \geq 4$ agents, and a finite set of payoff-relevant states Θ . Each agent i 's payoff is denoted $u_i(d, \theta)$, when a social decision $d \in D$ is selected in state θ .⁸ For example, in an auction, d is a vector of bidders' winning probabilities and their expected payments, and each u_i is quasi-linear in the payment part. Later, each agent's payoff *net* his information acquisition cost is considered as his objective.

At the beginning, neither the principal nor any of the agents know θ . The agents can acquire costly information about θ by generating private signals, possibly correlated to each other, whereas the principal cannot acquire any information about θ . Each agent has access to a sufficiently large set of possible signal realizations S_i . In principle, S_i (in particular, its size) may be a part of i 's choice, but assuming exogenous S_i is without loss of generality as long as $|S_i| \geq |\Theta|$.

We assume that information acquisition is fully private in the sense that neither the principal nor any other agent observes which information acquisition action i takes and which signal realization is observed by agent i . Agent i 's objective is the *net* payoff $u_i(d, \theta) - c_i(\sigma_i)$, where σ_i represents i 's information acquisition action. More about σ_i is discussed later. We assume the information acquisition cost function of agent i has the following properties:

Assumption 1. Properties of information acquisition cost.

1. $c_i(\sigma_i) \geq 0$ for any σ_i .
2. $c_i(\sigma_i) = 0$ if signal $s_i \in S_i$ generated by σ_i and θ are stochastically independent.

It is a quite general cost function specification, which includes many well-known cost functions as special cases: entropy-based cost functions, more general posterior-separable cost functions, and even some posterior-non-separable ones⁹. Observe that the second property makes

⁸We can endow the principal with his own payoff function $u_0(d, \theta)$, though it is not necessary.

⁹See the literature of cost of information, such as [Sims \(2003\)](#), [Matejka and McKay \(2015\)](#), [Caplin and Dean \(2015\)](#), and [Pomatto et al. \(2020\)](#).

sure that agent i pays nothing as long as he learns nothing from his signal s_i about the payoff-relevant state θ . The second assumption may either be seen as the “first-step” assumption, or be motivated by the competitive data provider story in the introduction. It may also be motivated by the following example. Suppose the payoff-relevant state is correlated with internet users’ online activity, and the agents are internet companies collecting users’ data. If an internet company breaches its users’ privacy, the users may legitimately claim compensation. The compensation will likely depend on the extent of the privacy breach, i.e. on the amount of information that is collected. Hence, if the internet company can demonstrate that its signal is independent of any individual user’s information, there is no breach of privacy, and thus, the internet company will not have to pay anything for acquiring an independent signal¹⁰.

So far, we discuss the properties of each agent’s information acquisition about θ and its cost. We now formalize the connection among their information acquisition actions. For example, suppose that agent i and j acquires information such that their signals s_i and s_j are imperfectly correlated with θ , in the form $s_i = \theta + \varepsilon_i$ and $s_j = \theta + \varepsilon_j$. How ε_i and ε_j should be correlated to each other? Most papers in the existing literature assume some specific correlation structures such as conditional (on θ) independence or perfect correlation. In this paper, we allow for flexible correlation structures, by *letting the agents choose the correlation structure as an equilibrium point*. For this purpose, we need a sufficiently large space of *the state of nature*, which is the source of all the random variables in the model ($\theta, \varepsilon_i, \varepsilon_j$, etc.), and let each agent’s information acquisition action be a choice of a probability distribution over this state of nature.

Formally, we assume that there is a space of the state of nature $X = [0, 1]$ with a typical element x , equipped with a Borel σ -algebra and a uniform probability measure \mathbb{P} .¹¹ Taking a richer space of the state of nature would not change our results. There is also a commonly known measurable function $\Theta : X \rightarrow \Theta$ mapping the state of nature to the payoff-relevant states. This

¹⁰In practice, major Internet companies (e.g. Google) claim that the way they collect (and use) information actually protect users’ privacy (e.g. according to the criterion of “differential privacy”), and in most cases they are able to collect information without paying anything to users.

¹¹See [Gentzkow and Kamenica \(2017\)](#) who adapt this approach in the context of multi-sender Bayesian persuasion.

function induces a common prior on the payoff-relevant states as follows: $\mu_0(\theta) \equiv \int_0^1 1_{\{\Theta(x)=\theta\}} dx$ for each $\theta \in \Theta$. Agent i 's information acquisition action is a measurable function $\sigma_i : X \rightarrow S_i$, such that, once x (and hence $\theta = \Theta(x)$) is realized, then i observes $s_i = \sigma_i(x)$. Let Σ_i denote the set of all such measurable functions, defining i 's information acquisition action space.

This way, for example, agent i can acquire $s_i = \theta + \varepsilon_i$ and agent j acquires $s_j = \theta + \varepsilon_j$ with any arbitrary correlation between ε_i and ε_j , and which correlation structure they choose is determined in an equilibrium. The cost function for σ_i is any function c_i that satisfies Assumption 1.

Any profile of information acquisition actions σ induces a joint distribution over payoff-relevant states and signal realizations, which we denote by $\alpha \in \Delta(\Theta \times S)$. When we want to make its dependence on σ more explicit, we write α_σ .

2.2 Mechanism

The principal faces both hidden action and hidden information of each agent. The principal commits to a mechanism at the *ex ante* stage in order to control the agents' incentives. More specifically, following the literature, we let the principal (i) send a message privately to each agent before his information acquisition action, and (ii) collect a message privately from each agent after the agent has observed a signal realization. Formally, a mechanism comprises $(R, \rho; M, \delta)$ where $R = (R_i)_{i=1}^I$ and $M = (M_i)_{i=1}^I$; R_i denotes the set of messages that the principal can send to each agent i ; M_i denotes the set of messages that each agent i can send to the principal; $\rho \in \Delta(R)$ is a distribution over the principal's messages, and $\delta : R \times M \rightarrow D$ denotes the decision rule.

The timing of the game is summarized as follows:

$t = 0$: $x \sim U(0, 1)$ is drawn but no one observes it.

$t = 1$: The principal designs a mechanism $(R, \rho; M, \delta)$.

$t = 2$: After observing the mechanism and receiving $r_i \in R_i$, each agent i privately chooses his information acquisition action $\sigma_i \in \Sigma_i$.

$t = 3$: Each agent i privately observes $s_i = \sigma_i(x)$, and privately sends $m_i \in M_i$ to the principal.

$t = 4$: The principal executes $d = \delta(r, m)$ where $m = (m_i)_{i=1}^I$.

Because no agent observes the other agents' actions or information (even noisily) at all, we consider Nash equilibrium as a solution concept. Then, applying the revelation principle of [Forges \(1986\)](#), we focus on *direct* mechanisms where (i) the principal directly recommends an information-acquisition action to each agent, and each agent directly reports a signal to the principal, and (ii) each agent finds it optimal to obey the recommended action and truthfully report his signal.¹²

Formally, a direct mechanism comprises $((\sigma_i)_{i=1}^I, (S_i)_{i=1}^I, \delta)$, where the principal recommends $\sigma_i \in \Sigma_i$ privately to each agent i ,¹³ and executes $\delta(s) \in D$ if the agents report $s = (s_i)_{i=1}^I \in S = \times_{i=1}^I S_i$. A direct mechanism is *incentive compatible* if it satisfies the following constraints: for any $\sigma'_i \in \Sigma_i$ and $\tau_i : S_i \rightarrow S_i$,

$$\sum_{\theta, s_i, s_{-i}} (u_i(\delta(s_i, s_{-i}), \theta) \alpha_{\sigma_i, \sigma_{-i}}(\theta, s_i, s_{-i})) - c_i(\sigma_i) \geq \sum_{\theta, s_i, s_{-i}} (u_i(\delta(\tau_i(s_i), s_{-i}), \theta) \alpha_{\sigma'_i, \sigma_{-i}}(\theta, s_i, s_{-i})) - c_i(\sigma'_i).$$

That is, each i must find it optimal to obey the recommended σ_i and report the realized s_i truthfully.

Although the constraints are concisely summarized by the inequalities above, they are actually rather complicated. First, changing σ_i affects the joint distribution α of (θ, s) and the agent's cost in a non-trivial way since agent i cannot affect agent $-i$'s information structure. Second, an agent may potentially want to make a double deviation, that is, change σ_i and at the same time change his reporting strategy.

¹²The proof proceeds as follows. First, imagine an auxiliary game where there is no principal, but instead, there is a fictitious player ("player 0") who is indifferent across all decisions in any state. At first, each agent i plays σ_i privately, and then observes the realized signal s_i privately. Then, (without any communication), player 0 chooses $d \in D$. Interpreting this as a baseline extensive-form game, it is easy to see that our current game (with the principal) is the mediated communication game of this auxiliary game in the sense of [Forges \(1986\)](#) (see also [Myerson \(1986\)](#)). Thus, her revelation principle applies.

¹³We focus on a deterministic recommendation of σ , rather than any stochastic recommendation. Accordingly, δ is denoted simply by $\delta(s)$ instead of $\delta(r, s)$. Since first best implementation is achieved with pure recommendations, our focus on them is without loss of generality.

Remark 1. Here, we do not explicitly impose individual rationality or participation constraints. It is not difficult to accommodate these constraints: let us require that any feasible direct mechanism must have an extra message m_i^\emptyset (a “non-participation” message) so that i ’s message space is now $S_i \cup \{m_i^\emptyset\}$, and $\delta(m_i^\emptyset, m_{-i})$ is some specific allocation (a “non-participation allocation”) for agent i , for any given m_{-i} . When the non-participation message is included into the set of messages for each agent, the individual rationality constraints, both at the *ex ante* and *ex interim* stages, are captured by the above incentive compatibility constraints.¹⁴

3 Main result

Fix any function $d^* : \Theta \rightarrow D$, which describes all the economically relevant outcomes in this environment except for the information acquisition costs. If the principal could observe θ , then any d^* is attainable without any information acquisition cost on the agents’ side. In this sense, one may interpret this d^* together with zero cost for the agents as the *first-best* outcome.¹⁵

In this section, for any given d^* , we explicitly construct a mechanism that implements d^* at zero cost for the agents. That is, the first best outcome can be attained *even though the principal cannot directly observe θ* .

Theorem 1. Fix any $d^* : \Theta \rightarrow D$. Under Assumption 1, there exists a mechanism (σ, S, δ) such that (i) $\sum_s \alpha(\theta, s) 1_{\{\delta(s) = d^*(\theta)\}} = \mu_0(\theta)$ for all θ , and (ii) $c_i(\sigma_i) = 0$ for all i .

Proof. The theorem is proved by construction.

Since Θ is a finite set, we assume without loss of generality that $\Theta = \{1, \dots, T\}$. Let $K > \max\{I, T\}$ be a prime number. Because \mathbb{P} is a uniform measure on $X = [0, 1]$, we can find a partition of X , denoted by $\{X_{\theta\psi}\}_{(\theta,\psi) \in \{1, \dots, T\} \times \{1, \dots, K\}}$, satisfying $\int_0^1 1_{\{x \in X_{\theta\psi}\}} dx = \frac{1}{K} \mu_0(\theta)$ for any θ and ψ . Define a measurable function $\Psi : [0, 1] \rightarrow \{1, \dots, K\}$ such that, if $x \in \cup_{\theta \in \Theta} X_{\theta\psi}$,

¹⁴Ex interim individual rationality is guaranteed because agent i can always deviate to $\tau_i(\cdot) \equiv m_i^\emptyset$. Ex ante individual rationality is guaranteed because agent i can always deviate to a costless σ_i and then to $\tau_i(\cdot) \equiv m_i^\emptyset$.

¹⁵For example, one may assume that $d^*(\theta)$ is the best decision of the principal given his own preferences in state θ .

then $\Psi(x) = \psi$. Immediately, Ψ is uniformly distributed on $\{1, \dots, K\}$ conditional on any realization θ of Θ , hence Ψ is independent of Θ .

Now consider the following information acquisition action profile: for each $i \in \{1, \dots, I\}$, $S_i = \{1, \dots, K\}$, and $\sigma_i(x) = \Theta(x) + i \cdot \Psi(x) \bmod K$ for any $x \in [0, 1]$. Note that the residual is calculated as in standard modular arithmetic except when $\Theta(x) + i \cdot \Psi(x)$ is divisible by K , in which case we set $\sigma_i(x) = K$ instead of 0. The following lemma gives the properties of (S, σ) that we need to prove the theorem.

Lemma 1. *The above (S, σ) satisfies:*

- (i) *For any $i \in \{1, \dots, I\}$, σ_i is independent of Θ .*
- (ii) *Conditional on any realization of (s_i, s_j) such that $i \neq j$, the joint distribution of Θ and $(\sigma_k)_{k \neq i, j}$ is degenerate.*

Proof of the lemma. By definition of (S, σ) , for each i , we have $s_i = \theta + i \cdot \psi \bmod K$, where the random variables Ψ and Θ are independent. Thus the signal profile $s = (s_i)_{i=1}^I$ is defined in the same way as in [Zhu \(2021\)](#).¹⁶ Thus, this lemma is directly implied by Lemma 2 in [Zhu \(2021\)](#). \square

The first property says that Θ and σ_i are independent, implying $c_i(\sigma_i) = 0$. The second property says that, given s_i, s_j with $i \neq j$, we can identify the true payoff-relevant state θ and any signal realization s_k without error, that is, there exist $\hat{\theta}(s_i, s_j)$ and $\hat{s}_k(s_i, s_j)$ such that:

$$\Pr(\Theta = \hat{\theta}(s_i, s_j) | s_i, s_j) = \Pr(\sigma_k = \hat{s}_k(s_i, s_j) | s_i, s_j) = 1.$$

Let the principal recommend the above σ , and offer the decision rule δ as follows: $\delta(s) = d^*(\theta)$ if (i) for any i, j with $i \neq j$, we have

$$\theta = \hat{\theta}(s_i, s_j);$$

or if (ii) there is i such that, for any j, k where i, j, k are all different, we have

$$\theta = \hat{\theta}(s_j, s_k).$$

¹⁶In fact, our signal profile s coincides with what [Zhu \(2021\)](#) calls “the IUAR disclosure policy, where IUAR is short for *individually uninformative but aggregately revealing*.”

In any other case, $\delta(s)$ is arbitrary.

Clearly, if the agents obey the recommendation and report their signals truthfully, then the first best outcome is attained. Therefore, we complete the proof by showing that the proposed mechanism satisfies incentive compatibility. Take any agent i , and suppose that he deviates to any σ'_i and reports $\tau_i(s_i)$ when s_i is realized. First, his cost of information acquisition increases weakly. Second, his reporting decision does not affect the social decision at all, because the principal executes $\delta(s) = d^*(\hat{\theta}(s_j, s_k))$ for an arbitrary pair (j, k) which does not include i . Therefore, the mechanism is incentive compatible. \square

4 Applications

4.1 Full-surplus extraction in common value auctions

Consider the following common value auction environment. The seller (principal) has a single indivisible good, and there are $I \geq 4$ bidders. The value of the good is common to all the bidders, denoted by $\theta \in \Theta$, where Θ is finite. In fact, the analysis of this section can be straightforwardly extended to the case of “non-pure” common values where each i ’s valuation is $v_i(\theta)$. Let $\mu_0(\theta)$ denote the probability that θ is the bidders’ common value.

Each bidder i ’s payoff is $\theta q_i - t_i - c_i(\sigma_i)$ if he wins the good with probability q_i , pays t_i to the seller, and spends $c_i(\sigma_i)$ as his information acquisition cost. In case he does not participate in the mechanism, his outside-option payoff is 0. The seller’s payoff is revenue, $\sum_{i=1}^I t_i$.

The first-best expected surplus of this society is the expected common value:

$$\sum_{\theta \in \Theta} \mu_0(\theta) \theta = \mathbb{E}[\theta].$$

There are several cases where the seller can easily earn $\mathbb{E}[\theta]$. First, if the seller *knows* θ , then he can simply post price θ . Even if the seller does not know θ , if the bidders know θ as their common knowledge (i.e., as *free* information), then again the seller can earn $\mathbb{E}[\theta]$. Conversely, if all the bidders are completely *uninformed* (so that each only knows the common prior μ_0), then again, the seller can post price $\mathbb{E}[\theta]$.

Notice that, with costly information acquisition as considered in our paper, neither of the above ideas would work. First, although it might be possible to make every bidder fully learn θ in some equilibrium, it does not yield $\mathbb{E}[\theta]$ as long as full information is strictly costly. Second, if the seller posts price $\mathbb{E}[\theta]$, then each bidder has a strong incentive of knowing whether the true θ is below $\mathbb{E}[\theta]$ or not: If i finds that $\mathbb{E}[\theta|s_i] < \mathbb{E}[\theta]$ given some signal s_i , he would not buy the good. As long as such information is not too costly, the bidder would be better off by acquiring it.

Therefore, with a general information acquisition cost function, the equilibrium information should be somewhere between full and no information, and it is *a priori* unclear how the seller should find the optimal balance of information and rent extraction. Nevertheless, as long as the cost functions satisfy Assumption 1, Theorem 1 implies that the full-surplus extraction is possible.

Corollary 1. *Under Assumption 1, there is a mechanism which yields $\mathbb{E}[\theta]$ as the seller's expected revenue (and each bidder earns 0).*

It is worth emphasizing that the logic here is very different from that of Crémer and McLean (1988). In their paper, the seller exploits an *exogenously given* correlated signal structure, in order to construct a side-bet scheme that extracts the entire surplus. In our case, each bidder can choose any information structure. Indeed, if he prefers, a bidder can choose an information structure such that his information is *independent* from all the other bidders' signals. The Crémer-McLean lottery scheme, therefore, does not work here. Also, in their mechanism, each bidder's payoff can be strictly negative *ex post*, while in our case, it is zero *ex post*. Indeed, if a mechanism offers a negative ex post payoff in our environment, a bidder would have a strong incentive of getting a signal indicative of that event and then report differently, in order to avoid such a negative payoff.

4.2 First-best implementation in collective decision-making

Consider a committee with a designer (principal) and $I \geq 4$ members (agents) deciding whether to hire or not to hire a job market candidate. Formally, $d \in D = \{h, nh\}$. The quality

of the candidate is $\theta \in \Theta$, which is unobserved *ex ante*. The designer and all members of the committee hold a common prior belief $\mu_0 \in \Delta(\Theta)$ about the candidate's quality.

The utility that each member obtains from hiring / not hiring the candidate is defined as follows:

$$u_i(d, \theta) = \begin{cases} u_i(\theta), & \text{if } d = h \\ 0, & \text{if } d = nh \end{cases}$$

Without loss of generality, we assume that $u_i(\theta) = k_i\theta$.

Only the committee members can acquire information about the candidate at cost $c_i(\sigma_i)$. The designer aims to maximize the expected sum of all members' gross utilities.¹⁷ That is, ideally, he wants to hire the candidate if and only if $\sum_i k_i\theta \geq 0$. The first best expected surplus of all committee members is given by:

$$\sum_{\theta | \sum_i k_i\theta \geq 0} \mu_0(\theta) \sum_i k_i\theta \equiv W^{FB}$$

It is useful to note that the existing literature (see e.g. [Li \(2001\)](#) and [Gerardi and Yariv \(2008\)](#)) typically assumes that the committee members have access to information structures whose realized signals are independently distributed across them, conditional on the state of the world. Under these restrictions, the first best outcome cannot be implemented. There are two main forces that prevent the committee from implementing the first best with these restricted information structures: free-riding problem and conflict of interest. First, when committee members have a conflict of interest, they may prefer not to report their own acquired information truthfully. Second, even if all members share a common preference, information could be underprovided relative to the social optimum, because it is essentially a public good used to make a collective decision. For example, [Li \(2001\)](#) suggests that distorting the decision rule away from the *ex post* optimal rule (which is optimal under exogenous information) could help to alleviate the free-riding issue.

In contrast to the previous literature, our results show that with more flexible (even though still costly) information acquisition the designer can implement the first best outcome. On

¹⁷The result extends to the case where the designer maximizes expected sum of members' net utilities (taking into account the information acquisition costs).

the one hand, having access to a wider range of information acquisition technologies enlarges the set of feasible deviations for the agents. On the other hand, the principal now has more flexibility in designing information structures recommended to the agents. Given these two opposing effects, it is not immediately clear *a priori* whether the first best outcome becomes more or less difficult to attain. It turns out that the second effect dominates: with a larger set of feasible mechanisms, the principal is able to incentivize the agents to acquire and report their information truthfully, no matter what social choice rule the principal is trying to implement. Indeed, Theorem 1 implies that the first best is implementable as long as Assumption 1 about the cost functions holds.

Corollary 2. *Under Assumption 1, there is a mechanism which yields W^{FB} as the total expected surplus (the decision is made under full information with no cost).*

Our construction helps to resolve both of the issues that prevent first best implementation with conditionally independent signals. Recall that it costs nothing for an agent to acquire an “*individually uninformative*” signal which is assigned to him under the optimal mechanism. Therefore, the distorted provision of information is resolved. Moreover, even if committee members have a conflict of interest, under our mechanism, they cannot do better than being truthful since any unilateral deviation can be detected by the designer.

5 Concluding remarks

It is quite natural that agents may desire to refine their information in response to the mechanism. This paper proposes one possible framework, based on a class of information acquisition cost functions, such that the cost of information depends on the informativeness of each agent’s signal about the state of the world, but on its informativeness about other agents’ signals. We show that such a specification leads to an extremely positive result. We conclude the paper with two remarks.

5.1 Interdependent cost functions

One natural criticism may be that our mechanism induces full information if the signals are aggregated, even though any single signal is completely uninformative: would it be reasonable to assume that such σ_i is costless? Because the answer is necessarily yes under Assumption 1, the question is essentially whether Assumption 1 itself is reasonable. Assumption 1 is satisfied, in particular, in any information acquisition environment, in which the cost of information acquisition does not depend on the correlation structure among signals but only depends on their individual informational content. We argued in Section 1 that there are information acquisition environments, for which this assumption is indeed a reasonable one. In general, however, the correlation structure might affect the cost of information acquisition in various ways. On the one hand, there seem to be cases where more positive correlation is more expensive. For example, fix agent 1’s private information, and consider agent 2. If acquiring a positively correlated information necessarily means that agent 2 must *steal* (perhaps a part of) agent 1’s information, more positive correlation will be more costly. On the other hand, there are opposite situations. [Strulovici \(2021\)](#), for example, considers the environment where hard evidence is scarce in the sense that, if one agent “picks up” a piece of evidence, then it becomes difficult for the others to get the same or similar evidence. On top of that, there also seem to be cases where *less correlated* signals are more costly: for example, imagine that there exist 2 newspapers, and σ_i corresponds to the decision of which newspapers to buy (including the options of buying both or none). Assume that the information in those newspapers is highly positively correlated, and assume that agent 1 buys (only) newspaper 1 for sure. In this case, agent 2 receives a positively correlated signal regardless of the newspaper he chooses, while in order to obtain something less correlated, he must buy *both* newspapers and orthogonalize newspaper 2’s information from newspaper 1’s information. Thus, less correlation is more costly.

This discussion suggests that we must think more about modeling the microstructure of information acquisition, in order to determine which correlation structures are more costly. Mechanism design with such more specific information acquisition cost structures would certainly be an interesting future direction, and we hope this article could serve as a first step

towards that direction.

5.2 Two or three agents

The individually-uninformative-but-aggregately-revealing signal structures require four or more agents to implement the first best at zero cost for the agents. With three agents, although it is possible to determine whether *some* agent has unilaterally deviated or not, it is not possible to identify who the deviator is (and hence not possible to identify the true θ). To see that, consider the following counterexample.

Counterexample 1. *Suppose that there are two payoff-relevant states, i.e. $\Theta = \{1, 2\}$ and consider the mechanism with $K = 5$. Computing $s_i = \theta + i\psi \pmod{5}$, we obtain:*

$\theta = 1$	Agent 1	Agent 2	Agent 3	$\theta = 2$	Agent 1	Agent 2	Agent 3
$\psi = 1$	2	3	4	$\psi = 1$	3	4	5
$\psi = 2$	3	5	2	$\psi = 2$	4	1	3
$\psi = 3$	4	2	5	$\psi = 3$	5	3	1
$\psi = 4$	5	4	3	$\psi = 4$	1	5	4
$\psi = 5$	1	1	1	$\psi = 5$	2	2	2

Suppose the principal observes an out-of-equilibrium signal realization profile $(2, 5, 4)$. There are two unilateral deviations that lead to this profile. First, the true profile might be $(2, 3, 4)$ in state $\theta = 1$ with agent 2 deviating. Second, the true profile might be $(1, 5, 4)$ in state $\theta = 2$ with agent 1 deviating. Hence, the principal cannot identify the deviator, nor can the principal infer the true state.

Note, however, that if there exists a social decision $d \in D$ that can serve as a severe punishment for all agents for any given θ , then out-of-equilibrium reports can be severely punished by the principal, and the same first best implementation result obtains.

With two agents, each agent has much more freedom. Indeed, the authors work on a separate project with two agents. There, even under Assumption 1, an extremely positive result similar to the one obtained in this paper does not generally hold. The optimal mechanism might involve

some costly information acquisition, and hence, the specification of the cost functions matters even more.

References

- BERGEMANN, D., B. BROOKS AND S. MORRIS, “Informationally Robust Optimal Auction Design,” *Working paper*, 2017.
- BERGEMANN, D. AND J. VÄLIMÄKI, “Information Acquisition and Efficient Mechanism Design,” *Econometrica* 70 (2002), 1007–1033.
- BIKHCHANDANI, S., “Information Acquisition and Full Surplus Extraction,” *Journal of Economic Theory* 145 (2010), 2282–2308.
- BIKHCHANDANI, S. AND I. OBARA, “Mechanism Design with Acquisition of Correlated Information,” *Economic Theory* 63 (2017), 783–812.
- BROOKS, B. AND S. DU, “Optimal Auction Design with Common Values: An Informationally-Robust Approach,” *Econometrica* 89 (2021), 1313–1360.
- CAPLIN, A. AND M. DEAN, “Revealed preference, rational inattention, and costly information acquisition,” *American Economic Review* 105 (2015), 2183–2203.
- CONDORELLI, D. AND B. SZENTES, “Information Design in the Holdup Problem,” *Journal of Political Economy* 128 (2020), 681–709.
- CRÉMER, J. AND R. MCLEAN, “Full Extraction of the Surplus in Bayesian and Dominant Strategy Auctions,” *Econometrica* (1988), 1247–1257.
- DENTI, T., “Unrestricted Information Acquisition,” *Working paper* (2020).
- FORGES, F., “An Approach to Communication Equilibria,” *Econometrica* 54 (1986), 705–731.
- GENTZKOW, M. AND E. KAMENICA, “Bayesian Persuasion with Multiple Senders and Rich Signal Spaces,” *Games and Economic Behavior* 104 (2017), 411–429.

- GERARDI, D. AND L. YARIV, “Information acquisition in committees,” *Games and Economic Behavior* 62 (2008), 436–459.
- GERSHKOV, A. AND B. SZENTES, “Optimal voting schemes with costly information acquisition,” *Journal of Economic Theory* 144 (2009), 36–68.
- GLEYZE, S. AND A. PERNOUD, “Informationally Simple Incentives,” *Working paper*, 2020.
- KALAI, A. T., E. KALAI, E. LEHRER AND D. SAMET, “A commitment folk theorem,” *Games and Economic Behavior* 69 (2010), 127–137.
- KRÄHMER, D., “Information Disclosure and Full Surplus Extraction in Mechanism Design,” *Journal of Economic Theory* 187 (2020).
- , “Information design and strategic communication,” *American Economic Review: Insights* 3 (2021), 51–66.
- LI, H., “A theory of conservatism,” *Journal of political Economy* 109 (2001), 617–636.
- LIU, Q., “Correlation and common priors in games with incomplete information,” *Journal of Economic Theory* 148 (2015), 49–75.
- MATEJKA, P. AND A. MCKAY, “Rational inattention to discrete choices: A new foundation for the multinomial logit,” *American Economic Review* 105 (2015), 272–298.
- MENSCH, J., “Screening Inattentive Agents,” *Working paper* (2020).
- MORRIS, S. AND M. YANG, “Coordination and Continuous Stochastic Choice,” *Working paper*, 2021.
- MYERSON, R., “Multistage Games with Communication,” *Econometrica* 54 (1986), 323–358.
- PERSICO, N., “Committee design with endogenous information,” *The Review of Economic Studies* 71 (2004), 165–191.

- PETERS, M. AND C. TRONCOSO-VALVERDE, “A folk theorem for competing mechanisms,” *Journal of Economic Theory* 148 (2013), 953–973.
- POMATTO, L., P. STRACK AND O. TAMUZ, “The Cost of Information,” *Working paper* (2020).
- RAVID, D., R. ANNE-KATRIN AND B. SZENTES, “Learning Before Trading: On the Inefficiency of Ignoring Free Information,” *Working paper*, 2020.
- RENAULT, J., L. RENOU, AND T. TOMALA, “Secure message transmission on directed networks,” *Games and Economic Behavior* 85 (2014), 1–18.
- RENOU, L. AND T. TOMALA, “Mechanism design and communication networks,” *Theoretical Economics* 7 (2012), 489–533.
- ROESLER, A.-K. AND B. SZENTES, “Buyer-Optimal Learning and Monopoly Pricing,” *American Economic Review* 107 (2017), 2072–2080.
- SHANNON, C., “Communication Theory of Secrecy Systems,” *Bell System Technical Journal* 28 (1949), 656–715.
- SHI, X., “Optimal auctions with information acquisition,” *Games and Economic Behavior* 74 (2012), 666–686.
- SIMS, C., “Implications of rational inattention,” *Journal of Monetary Economics* 50 (2003), 665–690.
- STRULOVICI, B., “Can Society Function Without Ethical Agents? An informational Perspective,” *Working paper* (2021).
- WILSON, R., “Competitive Bidding with Disparate Information,” *Management Science* 15 (1969), 446–448.
- YAMASHITA, T., “Revenue guarantee in auction with a (correlated) common prior and additional information,” *Working paper* (2018).

YANG, M., “Coordination with flexible information acquisition,” *Journal of Economic Theory* 158 (2015), 721–738.

ZHAO, X., “Heterogeneity and unanimity: Optimal committees with information acquisition,” *Working paper* (2016).

ZHU, S., “Private Disclosure with Multiple Agents,” *Working paper* (2021).